

Applied Data Science

Mod.8 – Neural Networks – day 1



CAMBRIDGE SPARK

@cambridgespark

Weekend outline

Today :

- ▶ Neural Nets basics (Multilayer Perceptron or MLPs)
- ▶ Convolutional Neural Networks (CNNs)

Tomorrow :

- ▶ CNNs for time series
- ▶ RNNs and LSTMs for classification and regression

Before starting...

A lot of notions to cover and some will require you to come back to it if you want to fully grasp the notions:

- ▶ **essentials:** how an MLP works, the issues, where NNs can work and where they don't, how to apply pre-trained neural nets.
- ▶ **intermediate:** convolution, pooling, batch-training, batch-normalisation, activation functions, input/output layers
- ▶ **advanced:** weight regularisation, step-size scheme, dropout, vanishing gradient

It's not so much the maths but rather NNs rely on *almost all* of classical ML → some concepts require time to process. We *strongly* recommend you work in pairs.

Neural Nets: the basics

► Introduction to NN

- *NN for classification*
- *The perceptron*

► Essential Maths

- *Activation functions and loss functions*
- *Gradient descent*


► Training and using a NN

Neural Net = set of interconnected neurons

- ▶ flexible class of models for supervised learning
- ▶ neuron = small *processing unit* transforming an input (stimulus) into a signal
- ▶ **aim:** imitate a complex function

Standard role for NN: **classification**

- ▶ **training set** = $\{(\text{object}, \text{class})_i\}_{i=1:N}$
- ▶ **goal**: label new objects
- ▶ **function to imitate**: classification function
- ▶ **Example**: classifying handwritten digits




The MNIST dataset

- ▶ **objects** : images of handwritten digits (28×28 px)
- ▶ **labels** : corresponding digit: $\{0, \dots, 9\}$.
- ▶ **dataset** : 70k labeled images,
(60k for training, 10k for testing)


Note: this dataset has been shown to be *too-simple* but it helps to gain intuition. You will practice on more complicated datasets later.

Perceptron = 1-neuron NN



- ▶ **input** x with d components
- ▶ **weight** w with d components
- ▶ **activation function** f with **bias** b
- ▶ **output** \hat{y} (class)

Inside a simple neuron...



- compute the **weighted sum** $s = \langle x, w \rangle$ (*stimulus*)

$$\langle x, w \rangle = \sum_{i=1}^d w_i x_i$$

- compare **stimulus** s to **bias** b (*activation*):

$s > b$: return $\hat{y} = 1$ (*neuron fires*)

$s \leq b$: return $\hat{y} = 0$ (*neuron does not fire*)

Coding a perceptron...


Head to your notebook and...

- ▶ see how a basic perceptron function can be coded
- ▶ try to use it to classify points
- ▶ what do you need to set? are there multiple solutions?

Perceptron: checkpoint 1

```
def outPerceptron(x, w, b):  
    innerProd = np.dot(x, w)  
    output = 0  
    if innerProd > b:  
        output = 1  
    return output  
  
def multiOutPerceptron2(X, w, b):  
    return (np.dot(X, w) > b).astype(float)
```

Perceptron: checkpoint 2



Coding a perceptron...

Key points...

- ▶ parameters to set: **weights** w and **bias** b
- ▶ finding good parameters = **training process**
- ▶ the neuron you coded answers a binary question:


"Does the input have a specific feature or not?"

Needed to train a NN: a **loss** and an **optimiser**

- ▶ a **loss function** is a way to measure the performances corresponding to a set of parameters $\theta = (w, b)$
- ▶ an **optimiser** is a procedure to go from a θ^0 to an *improved* θ^1 with lower loss.


Recall: this is done in the *training phase*, performance and loss are measured on the **training set**.

Activation function : a teaser



(more later...)

Neural network = set of **interconnected neurons**



- ▶ an **input layer** with d nodes
- ▶ **hidden layers** each with a fixed number of neurons
- ▶ an **interconnection pattern** between layers
- ▶ an **output layer** with K nodes

The **architecture** : a hard choice

- ▶ **Depth** : look at *complex interactions* of features
- ▶ **Interconnection pattern** : exploit problem-specific structure
- ▶ **Ideal model** :
 - not too expensive to train → *simple architecture*
 - not overfitting → *simple architecture*
 - capture complexity of data → *complex architecture*

The **architecture** : a hard choice

- ▶ for **generic** problems, no one knows
- ▶ for **specific** problems, some architectures are known to perform well, often **inspired** from our brain
 - e.g.: **convolutional NN** for image classification

(more about all this later...)

Playing with the architecture

One nice way to learn about the impact of architectural choices when dealing with Neural Networks is the *tensorflow playground*

<http://playground.tensorflow.org>


Training = setting the parameters

- ▶ As for the perceptron, get **good parameters** but for all neurons this time
- ▶ Typical architectures have thousands of neurons → **many parameters** to train (often millions)
- ▶ Loss functions defined over (typically) **huge training sets**
- ▶ The training of a NN is **computationally hard** but can be done using **a simple class of optimisers** .

ESSENTIAL MATHS


Sigmoid and hinge activation functions

- Sigmoid activation functions : smooth approximations of the step function
 - *tanh, logistic,...*




Sigmoid and hinge activation functions

- ▶ **Hinge activation functions** : idea of neuron firing more frequently when stimulated more
 - *Rectified Linear Units (ReLU), Leaky ReLU, softplus, ...*



The softmax for the last layer

- ▶ No parameters, converts the K outputs into K scores
- positive, sum to 1
- can be interpreted as probability of input being in that category



Loss functions : way to rank parameters

- ▶ The **misclassification loss** function: counts the number of instances inaccurately classified (problem: not smooth)
- ▶ The **quadratic loss** function: sum the squares of errors, (same loss used in standard regression):

$$L(\theta) = \sum_{i=1}^N (\sigma_i(\theta) - y_i)^2$$

- ▶ The **cross entropy loss** function:

$$L(\theta) = - \sum_{i=1}^N y_i \log \sigma_i(\theta) + (1 - y_i) \log(1 - \sigma_i(\theta))$$

Optimising the loss \rightarrow gradient descent

$$\theta^{k+1} = \theta^k - \delta^k \nabla L(\theta^k)$$

Recall:

- ▶ δ^k is the **step-size** at step k (standard schemes such as ADAM are popular for training NNs)
- ▶ $\nabla L(\theta^k)$ is the **gradient** of the loss at the previous step
- ▶ provably leads to minimiser **only** if the loss function is convex (not the case for NNs...)

Gradient of the loss or **backprop**

Objective function is a sum over datapoints (as usual):

$$L(\theta) = \sum_{i=1}^N L_i(\theta)$$

The gradient is therefore $\nabla L(\theta) = \sum_{i=1}^N \nabla L_i(\theta)$.

- ▶ how do we compute $\nabla L_i(\theta)$?
- ▶ that's potentially a big sum (usually N is in the millions) \rightarrow can we simplify this?

Computing $\nabla L_i(\theta)$

- ▶ the model is *layered* \rightarrow the loss function is a **composition**
- ▶ gradient of composition requires “the chain rule”
- ▶ **very efficient implementations exist** \rightarrow *central element in libraries such as TensorFlow or Torch.*

Dealing with a big sum

Instead of

$$\nabla L(\theta) = \sum_{i=1}^N \nabla_i(\theta)$$

use an **approximation** :

$$\hat{\nabla} L(\theta) = \frac{N}{|S|} \sum_{i \in S} \nabla L_i(\theta)$$

where S is a random subset of $\{1, \dots, N\}$, a **batch**.

This is the principle behind **stochastic gradient descent** (SGD).

Adam and Xavier

- ▶ The training of Neural Nets is done with SGD
- ▶ State-of-the-art step-size schemes take this into account for example:
 - [Xavier/Glorot Initialisation](#)
 - [Adam stepping scheme](#)

Brief explanation coming...

Adam stepping scheme (★★)

- keep track of recent gradient estimates in order to approximate the mean and the variance of the gradient at current step

$$\theta^{k+1} = \theta^k - \frac{\delta m^k}{\sqrt{\hat{\nu}^k} + \epsilon}$$

m^k and ν^k are estimates of the first and second moment of the gradients.

there are quite a few methods out there, no guarantees but empirical performances. ADAM is known to work quite well for Deep Learning.

Xavier initialisation (★★)

- ▶ generate initial weights from random normal with mean 0, variance $1/\tau$.
- ▶ initially, the neural network should not increase or decrease the variance of an instance that goes through it.
- ▶ variance of input = variance of output for a neuron
- ▶ Forward : requires $\tau = n_{\text{in}}$
- ▶ Backward : requires $\tau = n_{\text{out}}$

Compromise by Xavier Glorot & Joshua Bengio:

$$\tau = \frac{n_{\text{in}} + n_{\text{out}}}{2}$$

>>> Attacking MNIST!

Head to your notebook and...

- ▶ use `Keras` to write a 500x300 Neural Network
- ▶ train it and test it on MNIST

MNIST - checkpoint 1

```
model = Sequential()  
model.add(Dense(500, input_shape=(784,)))  
model.add(Activation('relu'))  
model.add(Dense(300))  
model.add(Activation('relu'))  
model.add(Dense(10))  
model.add(Activation('softmax'))
```


MNIST - checkpoint 2

```
model.compile(loss='categorical_crossentropy',  
              optimizer='adam', metrics=["accuracy"])  
  
model.fit(images_train, labels_train,  
          batch_size=100,  
          epochs=10,  
          verbose=2,  
          validation_data = (images_test, labels_test))
```

Convolutional Neural Networks

- ▶ Introduction to deep learning
- ▶ The convolution operator
- ▶ Convolutional Neural Networks (CNNs)
- ▶ Regularisation
- ▶ Inspecting VGG-16

Introduction to deep learning



Going deep with neural networks

- ▶ Simple fully-connected neural networks (as described already) typically fail on high-dimensional datasets (e.g. [images](#)).
 - Treating each pixel as an independent input...
 - ...results in $h \times w \times d$ new parameters per neuron in the first hidden layer...
 - ...quickly deteriorating as images become larger—requiring exponentially more data to properly fit those parameters!
- ▶ **Key idea:** [downsample](#) the image until it is small enough to be tackled by such a network!
 - Would ideally want to extract some useful features first...
- ▶ \implies exploit [spatial structure](#) !

The **convolution** operator




Enter the **convolution** operator

- ▶ Define a small (e.g. 3×3) matrix (the **kernel**, \mathbf{K}).
- ▶ Overlay it in all possible ways over the **input image**, \mathbf{I} .
- ▶ Record **sums of elementwise products** in a new image.


$$(\mathbf{I} * \mathbf{K})_{xy} = \sum_{i=1}^h \sum_{j=1}^w \mathbf{K}_{ij} \cdot \mathbf{I}_{x+i-1, y+j-1}$$

- ▶ This operator exploits **structure** —neighbouring pixels influence one another stronger than ones on opposite corners!


Convolution example




Convolution example




Convolution example




Convolution example




Convolution example




Convolution example




Convolution example




Convolution example



Convolution example (with padding)



Applying convolutions




Applying convolutions


- ▶ Just by observing the convolved image, can you tell what kind of **pattern** the kernel detects?
- ▶ How would you design a kernel that detects **vertical** edges?
- ▶ What would the following kernel **detect**?

```
kernel = np.array([[ 1,  1,  1],  
                   [ 0,  0,  0],  
                   [-1, -1, -1]])
```

Convolution with colour




Convolution with colour



Convolution with colour

- ▶ Can you design a filter to detect the **edge** of Grace Hopper's **left shoulder** ?
- ▶ [Hint: make sure the weights in your kernel add up to zero!]


Convolutional neural networks (**CNNs**)



Convolutional layer


- ▶ A convolutional layer is specified by several kernels, to be applied to the output image of the previous layer.
- ▶ Convolution with one of the kernels (and potentially applying an activation function to every pixel) provides a single **channel** of the output image.
- ▶ Start with random kernels—let the network learn optimal ones by itself!
 - **N.B.** all we're doing is multiplying inputs by weights and adding them together \implies we can learn in the same fashion as before!

Stacking convolutional layers




```
model.add(Conv2D(32, (3, 3),  
                activation='relu',  
                padding='same'))
```

Pooling layers




```
model.add(MaxPooling2D())
```

Stacking pooling layers




Putting it all together



```
model.add(Flatten())  
model.add(Dense(256))  
model.add(Activation('relu'))
```

Regularisation


Learning the sine function



Why now?

- ▶ With previously covered networks and problems, **overfitting** tends not to become an issue.
- ▶ However, with CNNs and most image recognition problems, this becomes an extremely major issue!
- ▶ We will cover two “**black magic**” methods that are extremely good in practice...


Dropout



- ▶ Randomly “kill” each neuron in a layer with probability p during training only...?!

```
model.add(Dropout(0.5))
```


Batch normalisation



```
model.add(BatchNormalization())
```

Batch normalisation

- Solution: **renormalise** outputs of the current layer across the current batch, $\mathcal{B} = \{x_1, \dots, x_m\}$ (but allow the network to “revert” if necessary)!

$$\mu_{\mathcal{B}} = \frac{1}{m} \sum_{i=1}^m x_i \quad \sigma_{\mathcal{B}}^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2$$
$$\hat{x}_i = \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \varepsilon}} \quad y_i = \gamma \hat{x}_i + \beta$$

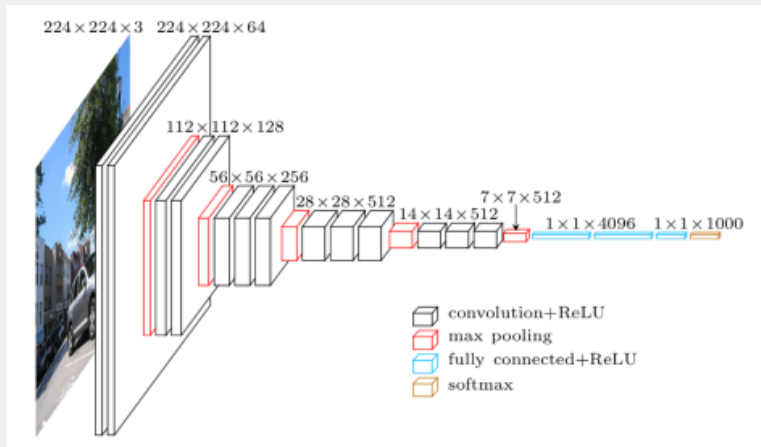
where γ and β are **trainable** !

- Now ubiquitously used across deeper CNNs
 - Published in February 2015, ~ 1600 citations by now!

One last trick: data augmentation

```
datagen = ImageDataGenerator(  
    width_shift_range=0.1,  
    height_shift_range=0.1)  
  
model.fit(...)  
model.fit_generator(datagen.flow(  
    X_train, y_train,  
    batch_size=32),  
    steps_per_epoch=len(X_train),  
    epochs=100,  
    validation_data=(X_test,  
        y_test))
```





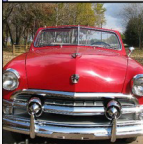

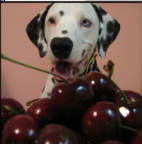
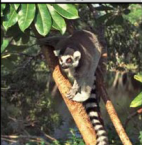
Inspecting VGG-16




ImageNet

- ▶ A 1000-class image classification problem, with classes that are both very **diverse** (animals, transportation, people...) and very **specific** (100 breeds of dogs!).
- ▶ A state-of-the-art predictor needs to be very good at extracting features from virtually any image!
- ▶ Early success story of deep learning (2012); human performance ($\sim 94\%$) surpassed by a 150-layer neural network in 2015.
- ▶ Pre-trained models are readily available in deep learning libraries (such as Keras, which I will be using for all the demos).



ImageNet classification

			
mite	container ship	motor scooter	leopard
<div> <div></div> <div>mite</div> <div>black widow</div> <div>cockroach</div> <div>tick</div> <div>starfish</div> </div>	<div> <div></div> <div>container ship</div> <div>lifeboat</div> <div>amphibian</div> <div>fireboat</div> <div>drilling platform</div> </div>	<div> <div></div> <div>motor scooter</div> <div>go-kart</div> <div>moped</div> <div>bumper car</div> <div>golfcart</div> </div>	<div> <div></div> <div>leopard</div> <div>jaguar</div> <div>cheetah</div> <div>snow leopard</div> <div>Egyptian cat</div> </div>
			
convertible	mushroom	cherry	Madagascar cat
<div> <div></div> <div>convertible</div> <div>grille</div> <div>pickup</div> <div>beach wagon</div> <div>fire engine</div> </div>	<div> <div></div> <div>agaric</div> <div>mushroom</div> <div>jelly fungus</div> <div>gill fungus</div> <div>dead-man's-fingers</div> </div>	<div> <div></div> <div>dalmatian</div> <div>grape</div> <div>elderberry</div> <div>ffordshire bullterrier</div> <div>currant</div> </div>	<div> <div></div> <div>squirrel monkey</div> <div>spider monkey</div> <div>titi</div> <div>indri</div> <div>howler monkey</div> </div>

Passing data through the network



Looking inside: the first layer



Looking inside: deeper layers

